




UNIVERSITÀ  
DEL SALENTO


# Clustering




# What is Cluster Analysis?




Given a set of objects, place them in groups such that the objects in **a group** are **similar** (or related) to one another and **different** from (or **unrelated** to) the objects in **other groups**



# Notion of a Cluster can be Ambiguous



How many clusters?



Six Clusters



Two Clusters



Four Clusters

# Types of Clusterings




A **clustering** is a set of clusters

Important distinction between **hierarchical** and **partitional** sets of clusters



- **Partitional Clustering:** A division of data objects into non-overlapping subsets (clusters)
- **Hierarchical clustering:** A set of nested clusters organized as a hierarchical tree

# Partitional Clustering

---




**Original Points**




**A Partitional Clustering**


# Hierarchical Clustering




**Traditional Hierarchical Clustering**



**Traditional Dendrogram**



**Non-traditional Hierarchical Clustering**



**Non-traditional Dendrogram**

# Other Distinctions Between Sets of Clusters

---



## **Exclusive** versus **non-exclusive** versus **Fuzzy**

- **Overlapping** or **non-exclusive** clusterings, points may **belong** to **multiple** clusters:
  - Can belong to multiple classes or could be 'border' points
- **Fuzzy** clustering (one type of non-exclusive)
  - In fuzzy clustering, a point **belongs to every cluster** with some weight between 0 and 1
  - Weights must sum to 1
- Probabilistic clustering has similar characteristics

## **Partial** versus **Complete**

- In some cases, we only want to cluster some of the data

# Types of Clusters

---



- a) Well-separated clusters
- b) Prototype-based clusters
- c) Contiguity-based clusters
- d) Density-based clusters
- e) Conceptual clusters



# Types of Clusters: Well-Separated




**Well-Separated Clusters:** A cluster is a set of points such that any point in a cluster is closer (or more similar) to every other point in the cluster than to any point not in the cluster.

Clusters that are quite far from each other.

*The distance between any two points in different groups is larger than the distance between any two points within a group.*

Well-separated clusters do **not need** to be **globular**, but can have any shape.




**3 well-separated clusters**

# Types of Clusters: Prototype-Based



**Prototype-based:** A cluster is a set of objects such that an object in a cluster is closer (more similar) to the prototype or "**center**" of a cluster, than to the center of any other cluster

The center of a cluster is often a **centroid**, the average of all the points in the cluster, or a **medoid**, the most "representative" point of a cluster







**4 center-based clusters**

# Types of Clusters: Contiguity-Based



**Contiguous Cluster** (Nearest neighbor or **Graph based**): A cluster is a set of points such that a point in a cluster is closer (or more similar) to one or more other points in the cluster than to any point not in the cluster.




**8 contiguous clusters**

# Types of Clusters: Density-Based



**Density-based:** A cluster is a dense region of points, which is separated by low-density regions, from other regions of high density.

Used when the clusters are irregular or intertwined, and when noise and outliers are present.



**6 density-based clusters**

# Types of Clusters: Conceptual cluster





**Conceptual cluster:** More generally, we can define a cluster as a set of objects that share some property.

This definition encompasses all the previous definitions of a cluster

New types of clusters --> clusters shown in figure: a triangular area is adjacent to a rectangular one

A clustering algorithm would need a **very specific concept** of a cluster to successfully detect these clusters.



# Characteristics of the Input Data Are Important

---



Type of proximity or density measure

- Central to clustering
- Depends on data and application

Data characteristics that affect proximity and/or density are

- Dimensionality
- Sparseness
- Attribute type
- Special relationships in the data (autocorrelation)

Noise and Outliers

# Clustering Algorithms

---



K-means

Hierarchical clustering

Density-based clustering

# K-means Clustering



**Partitional** clustering approach

Number of clusters, **K**, must be specified

Each cluster is associated with a **centroid** (center point)







Each **point** is assigned to the cluster with the **closest centroid**

The basic algorithm is very simple

- 
- 1: Select  $K$  points as the initial centroids.
  - 2: **repeat**
  - 3:   Form  $K$  clusters by assigning all points to the closest centroid.
  - 4:   Recompute the centroid of each cluster.
  - 5: **until** The centroids don't change
-



# Example of K-means Clustering



# K-means Clustering – Details

---



Simple iterative algorithm.

- Choose initial centroids;
- repeat {assign each point to a nearest centroid; re-compute cluster centroids}
- until centroids stop changing.

Initial centroids are often chosen randomly.

- Clusters produced can vary from one run to another

The centroid is (typically) the mean of the points in the cluster, but other definitions are possible

K-means will converge for common proximity measures with appropriately defined centroid

Most of the convergence happens in the first few iterations.

- Often the stopping condition is changed to 'Until relatively few points change clusters'

# K-means Objective Function



A common **objective function** (used with Euclidean distance measure) is


## Sum of Squared Error (SSE)

- For each point, the error is the distance to the nearest cluster center
- To get SSE, we square these errors and sum them.


$$SSE = \sum_{i=1}^K \sum_{x \in C_i} dist^2(m_i, x)$$

- $x$  is a data point in cluster  $C_i$  and  $m_i$  is the centroid (mean) for cluster  $C_i$
- SSE improves in each iteration of K-means until it reaches a local or global minima.


# Two different K-means Clusterings



**Original Points**

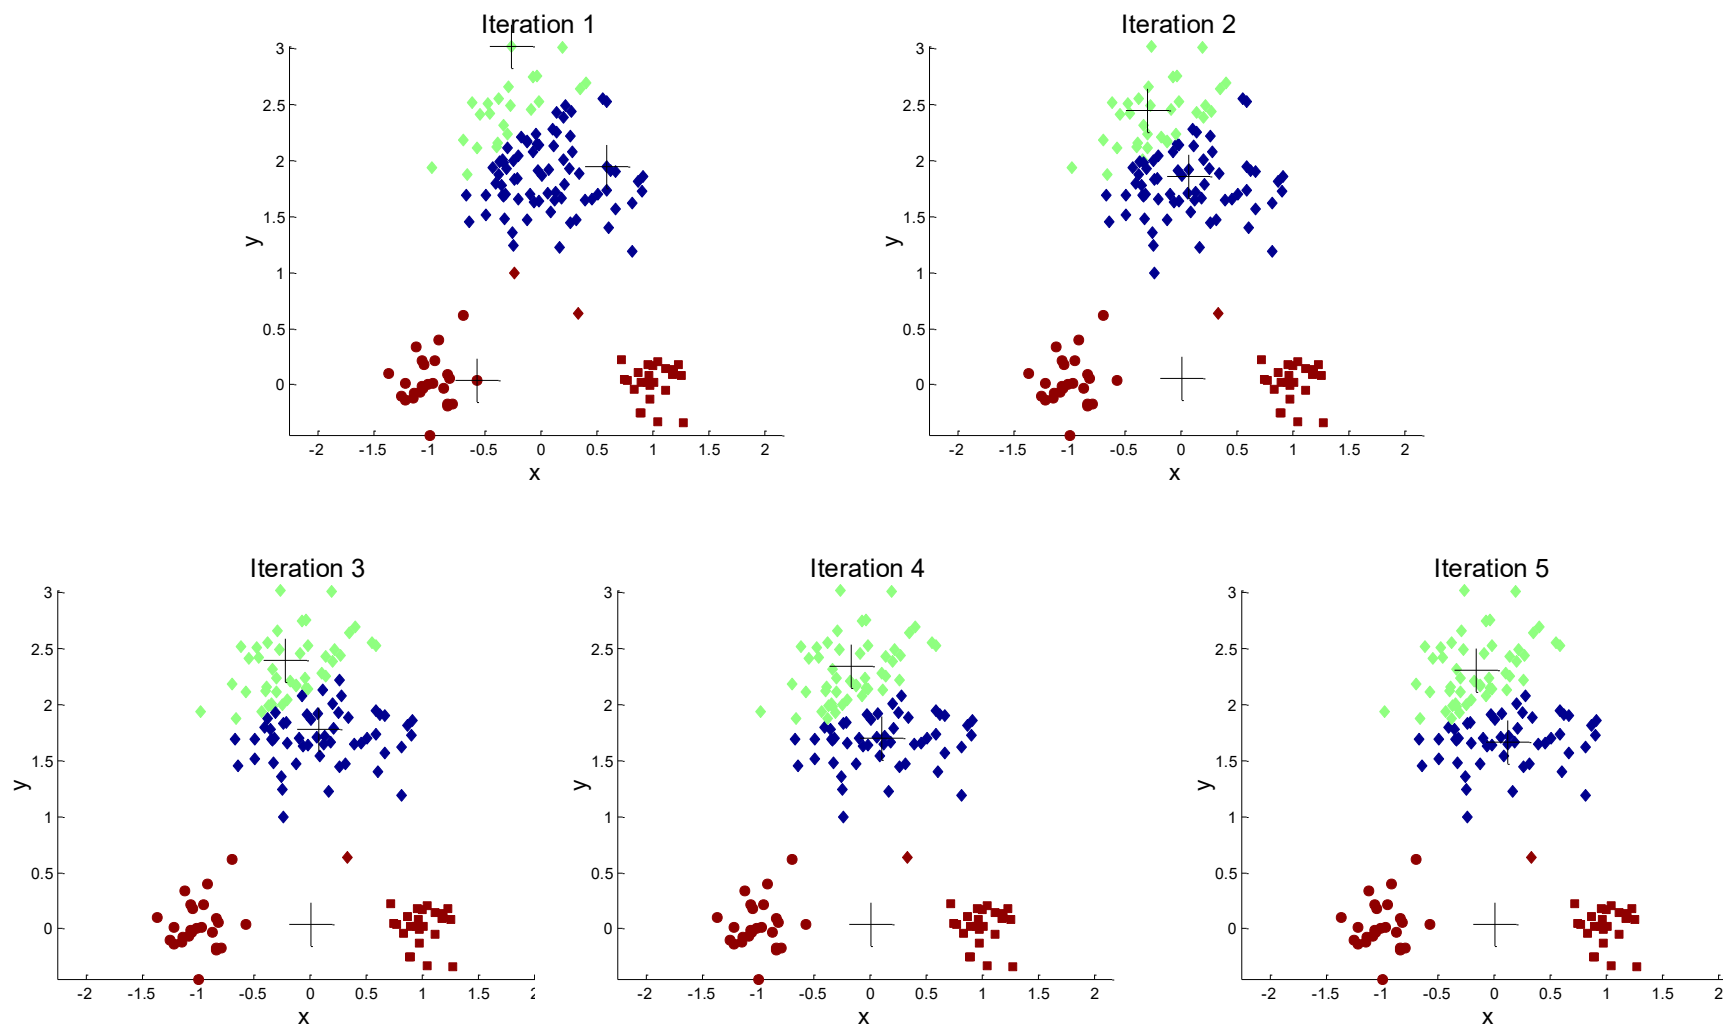


**Optimal Clustering**



**Sub-optimal Clustering**

# Importance of Choosing Initial Centroids ...



# Problems with Selecting Initial Points



If there are  $K$  'real' clusters then the chance of selecting one centroid from each cluster is small.

- Chance is relatively small when  $K$  is large
- If clusters are the same size,  $n$ , then


$$P = \frac{\text{number of ways to select one centroid from each cluster}}{\text{number of ways to select } K \text{ centroids}} = \frac{K!n^K}{(Kn)^K} = \frac{K!}{K^K}$$

- For example, if  $K = 10$ , then probability =  $10!/10^{10} = 0.00036$
- Sometimes the initial centroids will readjust themselves in 'right' way, and sometimes they don't
- Consider an example of five pairs of clusters



# Limits in random initialization: 10 Clusters Example

The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with two initial centroids in one cluster of each pair of clusters**



# Limits in random initialization: 10 Clusters Example

The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with two initial centroids in one cluster of each pair of clusters**





# Limits in random initialization: 10 Clusters Example

The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with two initial centroids in one cluster of each pair of clusters**

# Limits in random initialization: 10 Clusters Example



The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with two initial centroids in one cluster of each pair of clusters**

# Limits in random initialization: 10 Clusters Example



The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with some pairs of clusters having three initial centroids, while other have only one.**

# Limits in random initialization: 10 Clusters Example



The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with some pairs of clusters having three initial centroids, while other have only one.**

# Limits in random initialization: 10 Clusters Example



The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.




**Starting with some pairs of clusters having three initial centroids, while other have only one.**

# Limits in random initialization: 10 Clusters Example



The data consists of 5 pairs of clusters, where the clusters in each (top-bottom) pair are closer to each other than to the clusters in the other pair.



**Starting with some pairs of clusters having three initial centroids, while other have only one.**

# Solutions to Initial Centroids Problem

---



## Multiple runs

- Helps, but probability is not on your side

Use some strategy to select the  $k$  initial centroids and then select among these initial centroids

- Select most widely separated
  - K-means++ is a robust way of doing this selection
- Use hierarchical clustering to determine initial centroids

## Bisecting K-means

- Not as susceptible to initialization issues

# K-means++



This approach can be slower than random initialization, but very consistently produces better results in terms of SSE

To select a set of initial centroids,  $C$ , perform the following

1. Select an initial point at random to be the first centroid
2. For  $k - 1$  steps
3. For each of the  $N$  points,  $x_i$ ,  $1 \leq i \leq N$ , find the minimum squared distance to the currently selected centroids,  $C_1, \dots, C_j$ ,  $1 \leq j < k$ , i.e.,  $\min_j d^2(C_j, x_i)$
4. Randomly select a new centroid by choosing a point with probability proportional to  $\frac{\min_j d^2(C_j, x_i)}{\sum_i \min_j d^2(C_j, x_i)}$  is
5. End For



# Bisecting K-means




Variant of K-means that can produce a partitional or a hierarchical clustering


- 
- 1: Initialize the list of clusters to contain the cluster containing all points.
  - 2: **repeat**
  - 3:     Select a cluster from the list of clusters
  - 4:     **for**  $i = 1$  to *number\_of\_iterations* **do**
  - 5:         Bisect the selected cluster using basic K-means
  - 6:     **end for**
  - 7:     Add the two clusters from the bisection with the lowest SSE to the list of clusters.
  - 8: **until** Until the list of clusters contains  $K$  clusters
- 

**CLUTO:** <http://glaros.dtc.umn.edu/gkhome/cluto/cluto/overview>


# Bisecting K-means Example




# Bisecting K-means Example




# Bisecting K-means Example




# Bisecting K-means Example




# Bisecting K-means Example



# Bisecting K-means Example




# Bisecting K-means Example






# Bisecting K-means Example



# Bisecting K-means Example





# Limitations of K-means

---


K-means has problems when clusters are of differing

- Sizes
- Densities
- Non-globular shapes


K-means has problems when the data contains outliers.

- One possible solution is to remove outliers before clustering

# Limitations of K-means: Differing Sizes




**Original Points**




**K-means (3 Clusters)**

# Limitations of K-means: Differing Density




**Original Points**




**K-means (3 Clusters)**

# Limitations of K-means: Non-globular Shapes



**Original Points**




**K-means (2 Clusters)**

# Overcoming K-means Limitations




One solution is to find a **large number of clusters** such that each of them represents a part of a natural cluster.

Small clusters need to be put together in a **post-processing** step.



**Original Points**




**K-means Clusters**

# Overcoming K-means Limitations




One solution is to find a **large number of clusters** such that each of them represents a part of a natural cluster.

Small clusters need to be put together in a **post-processing** step.



**Original Points**



**K-means Clusters**




# Overcoming K-means Limitations




One solution is to find a **large number of clusters** such that each of them represents a part of a natural cluster.

Small clusters need to be put together in a **post-processing** step.



Original Points



K-means Clusters